

Introduction to reinforcement learning

Urtzi Ayesta and Matthieu Jonckheere

CNRS-IRIT & CNRS-LAAS

Markov reward processes and Markov decision processes

CHANGE OF NOTATIONS (we follow Barto and Sutton from now on).

Information - Set of policies

Markov Decision Process (MDP)

- ▶ observable state and reward
- ▶ known reward distribution and transition probabilities
- ▶ Action depends on current state and action, $p(s'|s, a)$

Partially Observable Markov Decision Process (POMDP)

- ▶ Partially observable state: we know z_t with known $P(s_t = s|z)$
- ▶ Observed rewards
- ▶ Known reward distribution and transition probabilities

Reinforcement learning

- ▶ observable state and reward
- ▶ Unknown reward distribution, unknown transition probabilities

Adversarial problems

- ▶ observable state and reward
- ▶ Arbitrary and time-varying reward function and state transitions

Classification of problems

Known model

MDP, POMDP

Unknown model

Reinforcement learning

Absence of model

Adversarial

Major components of an RL Agent

- ▶ Policy: A map from states to actions,
 $\pi(a|s) = \mathbb{P}(A_t = a|S_t = s)$
- ▶ Value function: Prediction of future rewards


$$v_\pi(s) = \mathbb{E}_\pi(R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s)$$

- ▶ A model predicts what the environment will do next

$$p(s'|s, a) = \mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a)$$

$$r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, A_t = a)$$

Maze Example




Rewards: -1 per time-step

Actions: N,E,S, W


States: Agent's location

Maze Example: Policy




Arrows represent policy $\pi(s)$

Maze Example: Value function



Numbers represent value $v_{\pi}(s)$ of each state s

Maze Example: Model




Dynamics: how actions change state

Rewards: From visited states
Model may be imperfect

Learning and planning

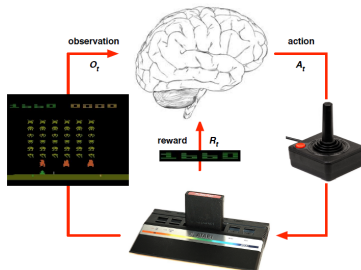
- ▶ Planning. A perfect model of the environment is known
⇒ MDP
- ▶ Reinforcement learning
 - ▶ Environment is initially unknown
 - ▶ Agent interacts with the environment
 - ▶ Agent learns policy

Atari Example: Planning




- ▶ Rules known
- ▶ Can query emulator
- ▶ If I take action a in state s , what would the score and next state be?
- ▶ Tree search to find optimal policy

Atari Example: Learning



- ▶ Rules unknown
- ▶ Learn directly from game-play
- ▶ Pick actions on joystick, see pixels and scores


Prediction and Control



a) gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

b) V_*



c) π_*

Markov Process

A Markov Process is a sequence of random states S_1, S_2, \dots , with the Markov property


Definition

A Markov Chain is a tuple $\langle \mathcal{S}, \mathcal{P} \rangle$


- ▶ \mathcal{S} is a (finite) set of states
- ▶ \mathcal{P} is a state transition probability matrix

$$p(s, s') = \mathbb{P}(S_{t+1} = s' | S_t = s)$$

Example: Student Markov Chain



Example: Student Markov Chain Transition Matrix



$$\mathcal{P} = \begin{matrix} & \begin{matrix} C1 & C2 & C3 & Pass & Pub & FB & Sleep \end{matrix} \\ \begin{matrix} C1 \\ C2 \\ C3 \\ Pass \\ Pub \\ FB \\ Sleep \end{matrix} & \begin{bmatrix} & & & & & 0.5 & \\ & 0.5 & & & & & 0.2 \\ & & 0.8 & & & & \\ & & & 0.6 & 0.4 & & \\ 0.2 & 0.4 & 0.4 & & & & 1.0 \\ 0.1 & & & & & 0.9 & \\ & & & & & & 1 \end{bmatrix} \end{bmatrix}$$

Markov Reward Process

A Markov Reward Process is a Markov Chain with values

Definition


Markov Reward Process is a tuple $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- ▶ \mathcal{S} is a (finite) set of states
- ▶ \mathcal{P} is a state transition probability matrix

$$p(s, s') = \mathbb{P}(S_{t+1} = s' | S_t = s)$$

- ▶ \mathcal{R} is a reward function $r(s) = \mathbb{E}(R_{t+1} | S_t = s)$
- ▶ $\gamma \in [0, 1]$ is a discount factor

Example: Student MRP



Markov Reward Process

A Markov Reward Process is a Markov Chain with values

Definition


Markov Reward Process is a tuple $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- ▶ \mathcal{S} is a (finite) set of states
- ▶ \mathcal{P} is a state transition probability matrix

$$p(s, s') = \mathbb{P}(S_{t+1} = s' | S_t = s)$$

- ▶ \mathcal{R} is a reward function $r(s) = \mathbb{E}(R_{t+1} | S_t = s)$
- ▶ $\gamma \in [0, 1]$ is a discount factor

Example: Student MRP



Return

Definition:

The return G_t is the total discounted reward from time-step t on:

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- ▶ The discount γ is the present value of future rewards
- ▶ The value of receiving reward R after $k + 1$ time-steps is $\gamma^k R$
- ▶ Immediate rewards more relevant than future ones
 - ▶ γ close to 0 referred as "myopic"
 - ▶ γ close to 1 referred as "far-sighted"

Why discount?

Most Markov reward and decision processes are discounted. Why?

- ▶ Mathematically convenient to discount rewards
- ▶ Avoids infinite returns in non-absorbing Markov processes
- ▶ If the reward is financial, immediate rewards may earn more interest than delayed rewards
- ▶ Animal/human behaviour shows preference for immediate reward

Value function


The value function $V(s)$ gives the long-run term value of state s

Definition:

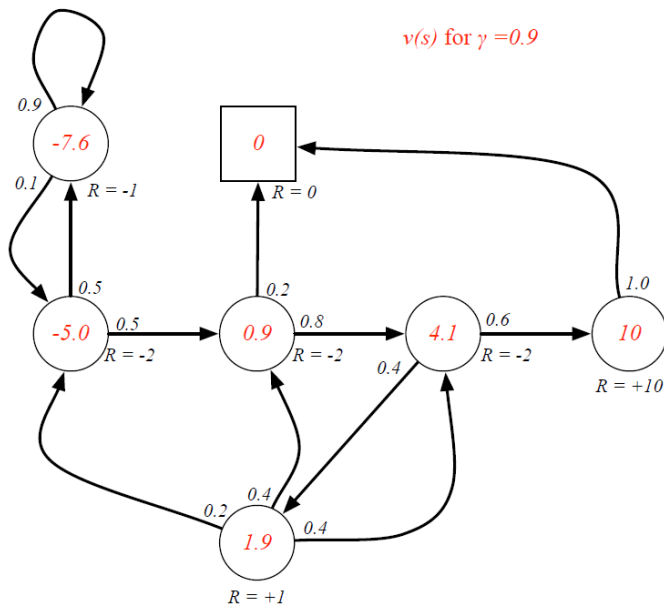
The state value function $V(s)$ of an MRP is the expected return starting from state s :

$$V(s) = \mathbb{E}(G_t | S_t = s)$$

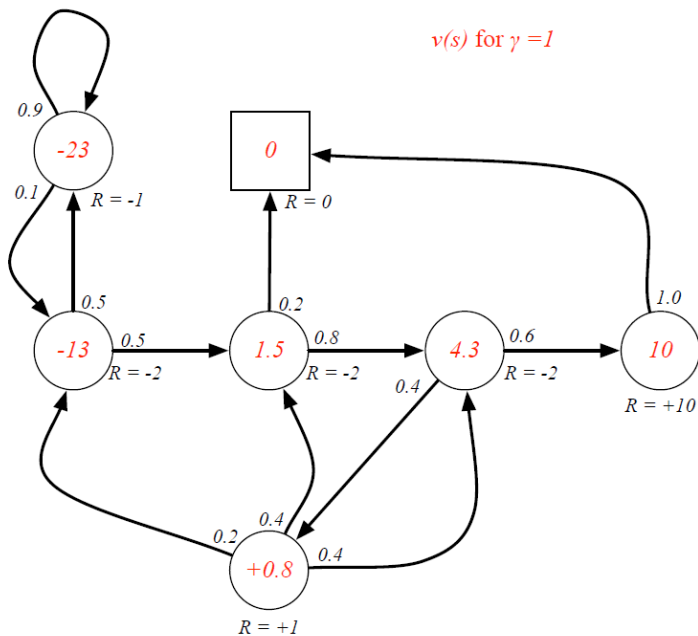
Example: Student MRP



Example: Student MRP



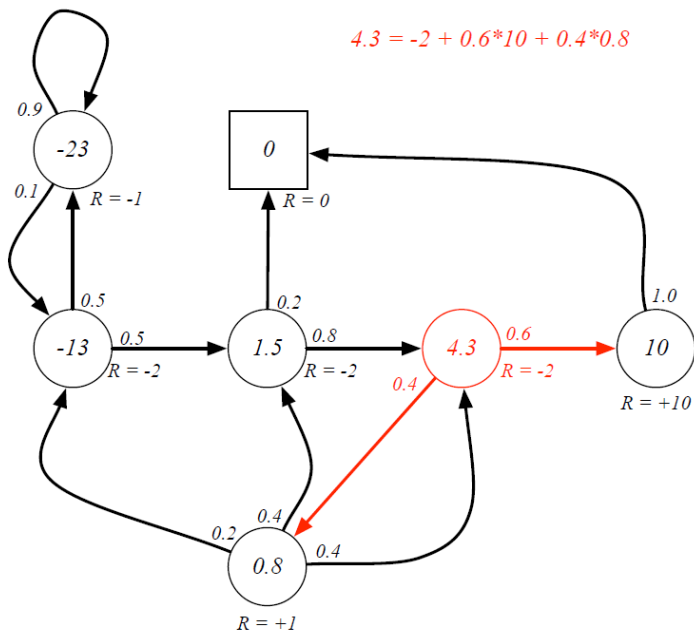
Example: Student MRP



Bellman Equation

$$\begin{aligned} V(s) &= \mathbb{E}(G_t | S_t = s) \\ &= \mathbb{E}(R_{t+1} + \gamma R_{t+2} + \dots | S_t = s) \\ &= r(s) + \gamma \sum_{s'} p(s, s') \mathbb{E}(R_{t+2} + \gamma R_{t+3} + \dots | S_{t+1} = s') \\ &= r(s) + \gamma \sum_{s'} p(s, s') V(s') \end{aligned}$$

Example: Student MRP



Bellman Equation in Matrix Form


$$\begin{pmatrix} V(1) \\ \vdots \\ V(n) \end{pmatrix} = \begin{pmatrix} r(1) \\ \vdots \\ r(n) \end{pmatrix} + \gamma \begin{pmatrix} p(1,1) & p(1,2) & \cdots & p(1,n) \\ p(2,1) & p(2,2) & \cdots & p(2,n) \\ \vdots & \vdots & \ddots & \vdots \\ p(n,1) & p(n,2) & \cdots & p(n,n) \end{pmatrix} \begin{pmatrix} V(1) \\ \vdots \\ V(n) \end{pmatrix}$$

It can be solved directly: $V = (I - \gamma\mathcal{P})^{-1}\mathcal{R}$

Computational complexity is $O(n^3)$ for n states

Many iterative methods for large MRP: dynamic programming,
Monte-Carlo evaluation, Temporal-Difference learning

Example: Grid World



Optimal policy when
 $R(s) = -0.04$ for every
non-terminal state

Policy Matrix: Each action is represented by a number : Action (Up) is represented by 0, (Rigth) by 1, (Down) by 2 and, finally, (Left) by 3

Example: Grid World (cont.)

Transition probabilities: Column 0 represents direction Up, Column 1 represents direction Right, Column 2 represents direction Down and Column 3 represents direction Left.

$$\begin{pmatrix} 0.8 & 0.1 & 0 & 0.1 \\ 0.1 & 0.8 & 0.1 & 0 \\ 0 & 0.1 & 0.8 & 0.1 \\ 0.1 & 0 & 0.1 & 0.8 \end{pmatrix}$$

Exercise: For $\gamma = 0.999$, calculate the value function for all the states

Markov Decision Process (MDP)

An MDP is a Markov Reward process with decisions.

Definition


Markov Reward Process is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- ▶ \mathcal{S} is a (finite) set of states
- ▶ \mathcal{A} is a finite set of actions
- ▶ \mathcal{P} is a state transition probability matrix

$$p(s'|s, a) = \mathbb{P}(S_{t+1} = s' | S_t = s, \mathcal{A}_t = a)$$

- ▶ \mathcal{R} is a reward function $r(s, a) = \mathbb{E}(R_{t+1} | S_t = s, \mathcal{A}_t = a)$
- ▶ $\gamma \in [0, 1]$ is a discount factor

Example: Student MDP



Policies (1)

Definition:

A policy π is a distribution over actions: :

$$\pi(s) = \mathbb{P}(A_t | S_t = s)$$

- ▶ A policy fully defines the behavior of an agent
- ▶ Policies are stationary or time-independent

Policies (2)

- ▶ Given an MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and a policy π
- ▶ The state sequence S_1, S_2, \dots is a Markov Process $\langle \mathcal{S}, \mathcal{P}^\pi \rangle$
- ▶ The state and reward sequence $S_1, R_2, S_2, R_3 \dots$ is a Markov reward process $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$

$$p^\pi(s, s') = \sum_{a \in \mathcal{A}} \pi(a|s) p(s'|s, a)$$

$$r^\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) r(s, a)$$

Value Function

Definition:

The state-value function $v_{\pi}(s)$ of an MDP is the expected return starting from state s , and then following policy π

$$V_{\pi}(s) = \mathbb{E}(G_t | S_t = s, A_{t:\infty} \sim \pi)$$


Definition:

The action-value function $q_{\pi}(s, a)$ is the expected return starting from state s , taking action a , and then following policy π

$$q_{\pi}(s, a) = \mathbb{E}(G_t | S_t = s, A_t = a, A_{t+1:\infty} \sim \pi)$$


Bellman Expectation Equation

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a)$$




Bellman Expectation Equation

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a)$$




$$q_{\pi}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_{\pi}(s')$$




Bellman Expectation Equation for V_π

$$V_\pi(s) = \sum_a \pi(a|s) \left(r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s') \right)$$




Bellman Expectation Equation for q_π

$$q_\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \sum_{a' \in \mathcal{A}} \pi(a'|s') q_\pi(s', a')$$



Bellman Expectation Equation in Student MDP



Optimal Value Function

Definition:

The optimal state-value function $V_*(s)$ is the maximum value function over all policies


$$V_*(s) = \max_{\pi} V_{\pi}(s)$$

The optimal action-value function $q_*(s)$ is the maximum action-value function over all policies


$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

An MDP is solved if we find either $V_*(s)$ or $q_*(s, a)$

Example: Optimal value function for Student MDP



Example: Optimal action-value function for Student MDP



Optimal Policy

Define a partial ordering over policies

$$\pi \geq \pi' \quad \text{if} \quad V_{\pi}(s) \geq V_{\pi'}(s), \forall s$$


Theorem:

For any Markov Decision Processes

- ▶ There exists an optimal policy π_* that is better or equal to all others, i.e., $\pi_* \geq \pi, \forall \pi$
- ▶ All optimal policies achieve the optimal value function $V_{\pi_*}(s) = V_*(s)$
- ▶ All optimal policies achieve the optimal action-value function $q_{\pi_*}(s, a) = q_*(s, a)$


Bellman Optimality Equation

$$V_*(s) = \max_{a \in \mathcal{A}} q_*(s, a)$$




Bellman Optimality Equation

$$V_*(s) = \max_{a \in \mathcal{A}} q_*(s, a)$$




$$q_*(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_*(s')$$




Bellman Optimality Equation for V_*

$$V_*(s) = \max_{a \in \mathcal{A}} \left(r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_*(s') \right)$$




Bellman Optimality Equation for q_*

$$q_*(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \max_{a' \in \mathcal{A}} q_*(s', a')$$



Bellman Optimality Equation in Student MDP




Finding an Optimal policy

An optimal policy can be found by maximizing over $q_*(s, a)$

$$\pi_*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_{a \in \mathcal{A}} q_*(s, a) \\ 0 & \text{otherwise} \end{cases}$$

Bellman Optimality Equation in Student MDP



Finite Markov Decision Process (MDP)

$$\begin{aligned} V_T^\pi(i) &= \mathbb{E}(R_1 + R_2 + \cdots + R_T | S_0 = i) \\ &= \mathbb{E}_\pi\left(\sum_{t=1}^T R_t | S_0 = i\right) \end{aligned}$$

$$V_t(i) = \sup_{\pi} V_t^\pi(i)$$

Finite Markov Decision Process (MDP)

Imagine action a , which yields reward $r(i, a)$. The best we can do from now on is


$$r(i, a) + \sum_j p(j|i, a) V^{T-1}(j)$$

Then, the best action is

$$V^T(i) = \max_a \left(r(i, a) + \sum_j p(j|i, a) V^{T-1}(j) \right)$$

Known as Optimality Equation, Dynamic Programming, Howard Equation,...

Bellman's key idea



Decision tree with depth T : $A^T S^{T+1}$ leaves (optimizing over history dependent policies)

Dynamic Programming: $S^2 A^T$ operations

Richard Bellman



1920 - 1984

American applied mathematician

Introduced **Dynamic Programming** (DP) as a method for solving a complex problem by breaking it down into a collection of simpler subproblems, solving each of those subproblems just once, and storing their solutions.

Extensions to MDPs

- ▶ Infinite and continuous MDPs
- ▶ Partially observable MDPs
- ▶ Undiscounted, average reward MDPs

Homework Assignment #1

Due date 23/04. To be done individually or by groups of 2.

- ▶ Using the equation in 31/61, derive the values for the value function of slides 26/61, 27/61 and 28/61. Explain the values obtained as γ changes.
- ▶ Verify that the values of the value function in slide 42/61 satisfy the Bellman Expectation Equation of 40/61 (as done in the example in red)
- ▶ Derive the values of the optimal action-value function q_* of 45/61 (Hint: To calculate $q(s, a)$ you can use the equation in 39/61, and then calculate q_* as in the definition)
- ▶ Verify and explain how the values in slide 50/61 satisfy the Bellman optimality equation.

Exercises

Exercise

A miner is at the bottom of a mine and sees three tunnels: 1, 2 and 3. Tunnel 1 leads to the exit in a hour. Tunnel 2 returns to the same crossroads in 2 hours and tunnel 3 returns to the crossroads in 3 hours. Every time the miner is at the crossroads, he chooses one of the tunnels with probability $1/3$, regardless of what he chose before. Define a Markov reward process describing this situation. Let T be the time it takes to leave the mine. Compute $E(T)$.

What would be the Bellman Optimality Equation ?

Exercise: Finite horizon MDP

Revenue management: Littlewood's model


Problem An airplane has 20 seats available, and the sell closes in 50 days. At every time epoch, the airplane decides the selling price: Either $p_1 = 5$, and then it will sell a seat with probability $q_1 = 0.1$, or $p_2 = 1$, and then it will sell a seat with probability $q_2 = 0.8$.

- ▶ Model the problem as a Finite Horizon MDP with total reward criterion, and write the optimality equation
- ▶ What is the optimal selling strategy? which qualitative conclusion you can draw from the solution?

Help: Let s denote the remaining seats available, and $V_T(s)$ denote the total reward in state s and with T days left. Use the principle of optimality of slide 54/61.

Exercise: Infinite Horizon MDP


Note: The reward from state 1 to 0 under action \ll is -15



What is the optimal policy (for total discounted reward) for various values of γ ?

- ▶ Write the optimality equation
- ▶ Solve it analytically as a function of γ

Exercise: "You're the reinforcement learner"



$$r(A, 1) = 10; r(A, 2) = -10; r(B, 1) = +40, r(B, 2) = +20$$

- Write the optimality equation