

# Introduction to reinforcement learning

Urtzi Ayesta and Matthieu Jonckheere

CNRS-IRIT & CNRS-LAAS

# Outline

- ▶ Principle of optimality
- ▶ Dynamic Programming
- ▶ Bellman Expectation as an Operator: Iterative Policy Evaluation
- ▶ Bellman Optimality Equation as an Operator: Value Iteration
- ▶ Policy Iteration

## Principle of optimality

**The tail of an optimal policy is optimal for the "tail" problem**

# Principle of optimality

**The tail of an optimal policy is optimal for the "tail" problem**

$$V_0^\pi(s) = \mathbb{E}_\pi\left(\sum_{t=1}^T R_t | S_0 = s\right)$$

Let  $\pi^* = (\pi_1, \dots, \pi_T)$  be the optimal policy. Then, the tail policy  $(\pi_k, \dots, \pi_T)$  is optimal for the tail cost

$$V_k^\pi(s) = \mathbb{E}_\pi\left(\sum_{t=k+1}^T R_t | S_k = s\right)$$

## Principle of optimality (cont.)

Imagine action  $a$ , which yields reward  $r(s, a)$ . The best we can do from now on is

$$r(s, a) + \sum_{s'} p(s'|s, a) V^{T-1}(s')$$

## Principle of optimality (cont.)

Imagine action  $a$ , which yields reward  $r(s, a)$ . The best we can do from now on is

$$r(s, a) + \sum_{s'} p(s'|s, a) V^{T-1}(s')$$

Then, the best action is

$$V^T(s) = \max_a \left( r(s, a) + \sum_{s'} p(s'|s, a) V^{T-1}(s') \right)$$

# What is Dynamic Programming

*Dynamic* sequential or temporal component to the problem

*Programming* : optimising a "program", i.e., a policy

- ▶ A method to solve complex problems by breaking them down into subproblems

# Planning by Dynamic Programming

Dynamic programming assumes full knowledge of MDP

It is used for planning in an MDP

For prediction:

- ▶ Input  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and policy  $\pi$
- ▶ output: value function  $V_\pi$

Or for control:

- ▶ Input  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
- ▶ output: optimal value function  $\pi_*$



## Bellman Expectation Equation for $V_\pi$

$$V_\pi(s) = \sum_a \pi(a|s) \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s') \right)$$

in case  $\pi(a|s) = 1, \forall s$ :

$$V_\pi(s) = \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s') \right)$$

## Bellman Optimality Equation for $V_*$

$$V_*(s) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_*(s') \right)$$

# Other Applications of Dynamic Programming

- ▶ Scheduling algorithms
- ▶ String algorithms (e.g. sequence alignment)
- ▶ Graph algorithms (e.g. shortest path algorithms)
- ▶ Graphical models (e.g. Viterbi algorithm)
- ▶ Bioinformatics (e.g. lattice models)

# Contraction Mapping Theorem

## Definition: $\gamma$ -contraction

We say that an operator  $T$  is a  $\gamma$ -contraction if

$$\|T(u) - T(v)\|_{\infty} \leq \gamma \|u - v\|_{\infty}$$

# Contraction Mapping Theorem

## Definition: $\gamma$ -contraction

We say that an operator  $T$  is a  $\gamma$ -contraction if

$$\|T(u) - T(v)\|_{\infty} \leq \gamma \|u - v\|_{\infty}$$

## Theorem: Contraction Mapping Theorem

Let  $\mathcal{V}$  be a complete metric space, en let  $T(\cdot)$  be a  $\gamma$ -contraction. Then:

- ▶  $T(\cdot)$  converges to a unique fixed point
- ▶ Convergence rate is exponential in  $\gamma$

# Bellman Expectation is a $\gamma$ -contraction

Define the *Bellman* expectation operator  $T^\pi$

$$T^\pi(V) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi V$$

$$T^\pi(V)(s) = \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right)$$

# Bellman Expectation is a $\gamma$ -contraction

Define the *Bellman* expectation operator  $T^\pi$

$$T^\pi(V) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi V$$

$$T^\pi(V)(s) = \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right)$$

$$\begin{aligned} \|T(U) - T(V)\|_\infty &= \|(\mathcal{R}^\pi + \gamma \mathcal{P}^\pi U) - (\mathcal{R}^\pi + \gamma \mathcal{P}^\pi V)\|_\infty \\ &= \|\gamma \mathcal{P}^\pi (U - V)\|_\infty \\ &\leq \|\gamma \mathcal{P}^\pi\| \|U - V\|_\infty \\ &\leq \gamma \|U - V\|_\infty \end{aligned}$$

# Iterative Policy Evaluation

*Problem* : Evaluate a given policy  $\pi$

- ▶ Solution: Iterations over Bellman equation
- ▶ At each iteration  $k + 1$ , update  $V_{k+1}(s)$  from state  $V_k(s')$

# Iterative Policy Evaluation

*Problem* : Evaluate a given policy  $\pi$


- ▶ Solution: Iterations over Bellman equation
- ▶ At each iteration  $k + 1$ , update  $V_{k+1}(s)$  from state  $V_k(s')$

Convergence of Iterative Policy Evaluation

- ▶ The Bellman expectation operator  $T^\pi$  has a unique fixed point
- ▶  $V_\pi$  is a fixed point of  $T^\pi$
- ▶ Iterative policy evaluation converges on  $V_\pi$



# Iterative Policy Evaluation



$$V_{k+1}(s) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_k(s')$$

# Bellman Optimality Equation

## Proposition:

$V^*$  is the unique solution of  $V^*$ :

$$V_*(s) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_*(s') \right)$$

Any stationary policy  $\pi^*$  that satisfies

$$\pi^*(s) \in \operatorname{argmax}_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_*(s') \right)$$

is an optimal policy

# Bellman Optimality Equation (proof sketch)

Consider the operator  $T_*$

$$T_*(V) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right)$$

and show that is a contraction mapping.

# Value Iteration

Starting with arbitrary  $V_0$ , define recursively  $\forall s$ :

$$V_{n+1}(s) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_n(s') \right)$$

Then  $\lim_{n \rightarrow \infty} V_n = V^*$ , where the rate of convergence is exponential

# Value Iteration

Starting with arbitrary  $V_0$ , define recursively  $\forall s$ :

$$V_{n+1}(s) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_n(s') \right)$$

Then  $\lim_{n \rightarrow \infty} V_n = V^*$ , where the rate of convergence is exponential

Retrieve the optimal policy  $\pi^*$  by:

$$\pi^*(s) \in \operatorname{argmax}_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_*(s') \right)$$

# Example: Shortest Path

g			

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

$V_1$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

$V_2$

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

$V_3$

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

$V_4$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

$V_5$


0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

$V_6$

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

$V_7$


# Evaluation a random policy in a small gridworld



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$   
on all transitions

# Evaluation a random policy in a small gridworld



- ▶ Undiscounted  $\gamma = 1$
- ▶ Nonterminal states  $1, \dots, 14$
- ▶ One terminal state (shown twice as shaded squares)
- ▶ Actions leading out of the grid leave state unchanged
- ▶ Reward is  $-1$  until the terminal state is reached
- ▶ Agent follows uniform random policy

$$\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 1/4$$



# Iterative Policy Evaluation in Small Gridworld

$v_k$  for the  
Random Policy

Greedy Policy  
w.r.t.  $v_k$

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	

random  
policy

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

	←	↕↕↕	↕↕↕
↑	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↓
↕↕↕	↕↕↕	→	

$k = 2$


0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

	←	←	↕↕↕
↑	↖	↕↕↕	↓
↑	↕↕↕	↘	↓
↕↕↕	→	→	

# Iterative Policy Evaluation in Small Gridworld (2)


$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0



$k = 10$


0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0



optimal policy

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



# How to improve a policy: Policy Iteration

Given a policy  $\pi$ , let  $V^\pi$  be its value function. We define the  $\bar{\pi} = \text{greedy}(V^\pi)$  to be the greedy policy with respect to  $V^\pi$

$$\bar{\pi}(s) \in \operatorname{argmax}_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s') \right) \quad \forall s.$$

# How to improve a policy: Policy Iteration

Given a policy  $\pi$ , let  $V^\pi$  be its value function. We define the  $\bar{\pi} = \text{greedy}(V^\pi)$  to be the greedy policy with respect to  $V^\pi$

$$\bar{\pi}(s) \in \operatorname{argmax}_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s') \right) \quad \forall s.$$

## Proposition: Policy Improvement

$V^{\bar{\pi}}(s) \geq V^\pi(s)$  and equality holds if and only if  $\pi$  is optimal

## Sketch of proof

Recall the operator  $T_*$  as

$$T_*(V) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right)$$

$$V^\pi = T_\pi V^\pi \leq T_*(V^\pi) = T_{\bar{\pi}}(V^\pi)$$

## Sketch of proof

Recall the operator  $T_*$  as


$$T_*(V) = \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right)$$

$$V^\pi = T_\pi V^\pi \leq T_*(V^\pi) = T_{\bar{\pi}}(V^\pi)$$

It can be shown that  $T_{\bar{\pi}}$  is a monotone operator, i.e.,  $V_1 \leq V_2$  implies  $T_{\bar{\pi}}(V_1) \leq T_{\bar{\pi}}(V_2)$ . Repeatedly applying  $T_{\bar{\pi}}$  to the above inequality


$$V^\pi \leq T_{\bar{\pi}} V^\pi \leq (T_{\bar{\pi}})^2 V^\pi \leq \dots \leq \lim_{n \rightarrow \infty} (T_{\bar{\pi}})^n V^\pi = V^{\bar{\pi}}$$

# Policy Iteration Algorithm



**Policy evaluation** Estimate  $v_\pi$   
 Iterative policy evaluation

**Policy improvement** Generate  $\pi' \geq \pi$   
 Greedy policy improvement



# Policy Iteration Algorithm (cont.)

## Pseudocode

- ▶ 0. Initialization: choose a policy  $\pi_0$

For  $k = 0, 1, \dots$ :

- ▶ Policy Evaluation: Compute  $V_{\pi_k}$  for example

$$V_{\pi_k} = (I - \gamma \mathcal{P}^{\pi_k})^{-1} \mathcal{R}^{\pi_k}$$

- ▶ Policy improvement: Compute  $\pi_{k+1} = \text{greedy}(V_{\pi_k})$
- ▶ Stop if  $V_{\pi_{k+1}} = V_{\pi_k}$ , else repeat

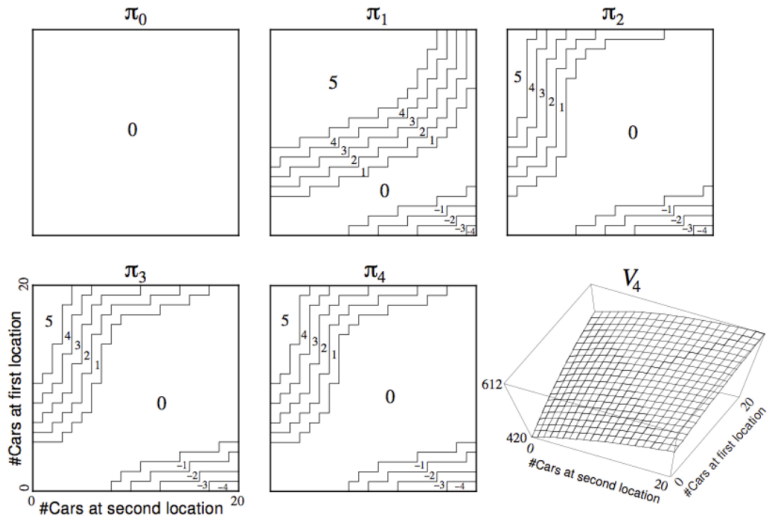


# Car's rental



- ▶ States: Two locations, maximum of 20 cars at each
- ▶ Actions: Move up to 5 cars between locations overnight
- ▶ Reward: 10 for each car rented (must be available)
- ▶ Transitions: Cars returned and requested randomly
  - ▶ Poisson distribution,  $n$  returns/requests with prob  $p(n)$
  - ▶ 1st location: average requests = 3, average returns = 3
  - ▶ 2nd location: average requests = 4, average returns = 2

# Cars's rental (cont.)



# Synchronous Dynamic Programming Algorithm

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iter. PI
Control	Bellman Expectation Equation + Greedy PI	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- ▶ Algorithms are based on state-value function  $V_{\pi}(s)$  or  $V_{*}(s)$
- ▶ Complexity  $O(mn^2)$  per iteration, for  $m$  actions and  $n$  states
- ▶ Could also apply to action-value function  $q_{\pi}(s, a)$  or  $q_{*}(s, a)$ , complexity  $O(m^2n^2)$  per iteration

# Asynchronous Dynamic Programming

- ▶ Synchronous backups  $\implies$  all states are backed up in parallel
- ▶ Asynchronous DP backs up states individually, in any order
- ▶ For each selected state, apply the appropriate backup
- ▶ Can significantly reduce computation
- ▶ Guaranteed to converge if all states continue to be selected

# Asynchronous Dynamic Programming

Three simple approaches

- ▶ In-place dynamic programming
- ▶ Prioritised sweeping
- ▶ Real-time DP

# In-place dynamic programming

In-place value iteration only stores one copy of value function.

$$V(S) \leftarrow \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right)$$

# Prioritized Sweeping

Use magnitude of Bellman error to guide state selection, e.g

$$\left| \max_{a \in \mathcal{A}} \left( r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right) - V(s) \right|$$

- ▶ Backup the state with the largest remaining Bellman error
- ▶ Update Bellman error of affected states after each backup
- ▶ Requires knowledge of reverse dynamics (predecessor states)
- ▶ Can be implemented efficiently by maintaining a priority queue

# Real-Time Dynamic Programming


- ▶ Idea: only states that are relevant to agent
- ▶ Use agent's experience to guide the selection of states
- ▶ After each time-step  $S_t, A_t, R_{t+1}$
- ▶ Backup the state  $S_t$

$$V(s_t) \leftarrow \max_{a \in \mathcal{A}} \left( r(s_t, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s_t, a) V(s') \right)$$



# Full-width backups

- ▶ DP uses full-width backups
- ▶ For each backup (sync or async)
  - ▶ Every successor state and action is considered
  - ▶ Using knowledge of the MDP transitions and reward function
- ▶ DP is effective for medium-sized problems (millions of states)
- ▶ For large problems DP suffers Bellman's curse of dimensionality
  - ▶ Number of states  $n = |\mathcal{S}|$  grows exponentially with number of state variables
- ▶ Even one backup can be too expensive



## Average Reward MDP

An ergodic MDP has an average reward per time-step  $g^\pi$  that is independent of start state


$$g^\pi = \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E} \left[ \sum_{t=1}^T R_t \right]$$

We define the value function as

$$\begin{aligned} V_\pi(s) &= \mathbb{E}_\pi \left[ \sum_{k=1}^{\infty} (R_{t+k} - g^\pi) | S_t = s \right] \\ &= \mathbb{E}_\pi [(R_{t+1} - g^\pi) + V_\pi(S_{t+1}) | S_t = s] \\ &= r(s, a) - g^\pi + \sum_{s'} p(s' | s, a) V_\pi(s') \end{aligned}$$

## Exercise: Infinite Horizon MDP

Note: The reward from state 1 to 0 under action  $\ll$  is  $-15$




What is the optimal policy (for total discounted reward) for various values of  $\gamma$ ?

- ▶ Solve the optimality equation it by value iteration
- ▶ Characterize the optimal policy using policy iteration

**Help:** You might use Value Iteration (slide 15/35), or policy iteration (slide 23/35)

## Exercise: Grid World



Optimal policy when  
 $R(s) = -0.04$  for every  
non-terminal state

Find the optimal policy using the Policy Iteration Algorithm (slide [32/45](#)).

Deadline: 14 Mai